

Sr. NO	Asked by	To be answered in the session live?	Category	Question	Answer
1	Bheem Singh		PO	Are all participants in today's orientation session new, meaning they have joined the program last week or so and this is the first session.	This is the program orientation session i.e. official launch for all learners enrolled in the January Cohort. The week 1 of the program starts next week.
2	Abhishek Trehan		PO	Hoping we will get recordings for all sessions incase we aren't able to join or want to replay for understanding	Absolutely! All the live sessions are recorded. The recordings are available on Olympus portal under the Activities tab within 6 hours of the session.
3	Soumyadip Mukhopadhyay		PO	Hi..how many total enrolled for the program, and what will be individual cohort sizes?	We have close to 350 learners in the January Cohort. The cohort will be split in groups of 20-25 learners for your mentored learning sessions.
4	Abhishek Trehan		PO	Great and till when are these saved? I mean when do you revoke access?	The recordings along with the course content will be available for 3 years from enrollment.
5	Ram Krishnamurthy		AO	Will we be covering MCP in this course?	Yes, we will be covering the MCP in this program
6	Rodney Townsend		PO	Is there a way to get a calendar feed? Receiving many emails overall and many from great learning. Would love to be able to subscribe to events to easily get the join links and ensure I do not miss any sessions.	The Learning Schedule for the program with the relevant dates for all sessions, content release and assessments deadlines are available under the Program Overview Module under Courses on Olympus.
7	Tony Shan		PO	I looked at "All" learning activities on Olympus, and it turned out that only 3 items are shown, including the current orientation session. Can you post the calendar of all "live" sessions (required and optional) in this program, so I can plan ahead to block myself in my calendar as these time slots are on weekends? Please share the direct link or navigation path on the dashboard, if this info is available somewhere else on Olympus.	The Learning Schedule for the program with the relevant dates for all sessions, content release and assessments deadlines are available under the Program Overview Module under Courses on Olympus.
8	Dino Vitale		PO	Can I be put in cohort group with my co-worker? (Anmol Sharma)	Your program manager will connect with you regarding MLS groups alignment.
11	David Bernal		AO	Some people in my company say that agents only work in narrow use cases, is that the case? is design issue or just the state of the technology?	No, AI agents excel beyond narrow use cases across sales, supply chain, HR, and finance. Current limitations branches more from design issues like poor error handling, legacy integration, and orchestration flaws.
12	Meghan Osborn		PO	Are recording transcripts available inside the app? I didn't see them. If not, is it possible to get them added?	We do not have the transcripts available for the recorded lectures or live sessions on your end. However, you may reach out to your Program Manager, who can assist you with transcripts for each Mentored Learning Session.
13	Rodney Townsend		PO	I know the sessions are added to the dashboard, but it would be great to subscribe to a shared calendar. Other courses I have taken have this capability to ensure nothing is missed. I now know I have a session added tomorrow, but I was not aware until I logged in for today's session.	Regarding the recorded lectures and content, we have in-built LLM model, Glaide aligned with each lecture. You can use Glaide to ask questions or get a summary of any Video at your convenience.
14	Omotolani Lewis		PO	I have done and completed the pre-course learning material on Python. As someone without significant prior knowledge of coding, I am concerned that this may limit my performance on this course and how much I can really digest and get from it. How much of the course material relies on coding. Is the material from the precourse material sufficient grounding.	The assigned group and timing for your upcoming Mentored Learning sessions starting week 1 of the program will be shared by Wednesday next week. All the mentored learning sessions throughout the 16 learning weeks, will be at the same time.
15	Shankar Thangavelu		PO	Can we get a copy of the slideware as well along with the recordings ?	This is a coding-intensive program that involves hands-on work in Python throughout. The pre-work is specifically designed to help you build a strong foundation in Python, especially since you are starting without prior Python experience. Going through the pre-work diligently will ensure you are comfortable with the basics before the program begins. Additionally, Week 1 of the program serves as a refresher on Python for Agentic AI, giving you further confidence before moving into more advanced concepts.
16	David Bernal		AO	Will the course cover A2A protocol as well ?	If you follow the pre-work and Week 1 content thoroughly, you will be fully prepared to make the most of the program. On average, learners are required to commit 8-10 hours per week to cover the program content and attend weekend sessions. Since you are starting without prior Python experience, you may need to dedicate additional time to practice and revise Python concepts to stay confident and make the most of the hands-on components.
17	Ankur Vohra		PO	The program schedule shows only two master class from JHU I was told this will be refreshed in Jan at the start of program for more classes, by when this is expected ? Will there be quick walkthrough of schedule as well today ?	The Orientation slide deck will be available under Courses -> Program Overview and recording of the session will be found on Olympus under Completed Activities.
18	Ardendu Mishra		AO	Hi, what all tools/software we need to install and use on our system, for this program ?	Currently in the program we are only covering Model Context Protocol (MCP).
20	Bereket Haile		PO	The programme schedule only shows the dates not the time of the lessons. Can you please confirm the exact time this lessons will be?	We will have 4 masterclasses in the cohort. We will be adding the dates of the other 2 masterclasses in the learning schedule soon. Yes, there will be a quick walkthrough of the schedule towards the end of the session.
21	Meghan Osborn	LIVE	AO	"80% of AI initiatives fail..." Is this failure to launch or failure to succeed, or both? Curious if there are recurring or common friction themes.	You can install the VSCode + Python in your system or the alternative is to use the Google Colab
22	Roman Lozano		PO	What is the start date for week 1 of the program?	We will be able to share the timing for all your upcoming mentored learning sessions by Wednesday next week.
23	Ardendu Mishra		AO	Are we covering Langchain/Langgraph concepts in this program ?	It usually refers to both failure to launch and failure to achieve sustained value in production. Common friction points include unclear use case definition and success metrics, poor data quality and integration into enterprise systems, and weak evaluation that misses real failure modes. Teams also underestimate governance, security, and change management, and overestimate model reliability under distribution shift. Cost and latency surprises often appear after pilots when usage scales. Next, examine LLM Ops practices: dataset and prompt versioning, automated regression evals, and observability with incident response.
24	David Bernal		AO	So you are saying that 100% of the root causes are "human related" and not technology related?	The week 1 of the program starts on Jan 12th, your first weekend Mentored Learning Session will be on the weekend of Jan 17th/18th.
25	Manoj Kumar	LIVE	AO	What are your thoughts on the ethical, privacy, and security implications of AI's evolution, and how can it be controlled?	Yes, we cover LangChain and LangGraph in the program.
26	Anonymous Attendee	LIVE	AO	I have done and completed the pre-course learning material on Python. As someone without significant prior knowledge of coding, I am concerned that this may limit my performance on this course and how much I can really digest and get from it. How much of the course material relies on coding. Is the material from the precourse material sufficient grounding.	No, technical limits are real, but many projects fail because human decisions about scope, evaluation, integration, and governance determine whether those technical limits are managed or amplified.
27	Jeff O'Connel		AO	I'm a graduate of JHU Certificate in Generative AI. How much overlap of the content covered in the Gen AI curriculum is covered in this Agentic AI program?	Ethics, governance, guardrails, and compliance are major unresolved challenges in agentic AI. This topic is extensive and important enough to warrant a full dedicated module in the course. Governance is considered one of the biggest challenges in agentic AI today, especially as agents gain autonomy and privileges.
28	Aalpen Patel		AO	What is the best IDE, Vs code, Jupyter or Colab	No, you will have enough support from the team to progress through the course. The weekend mentor-led sessions are code-focused, but you can follow along and ask for help whenever you get stuck.
29	Avishek Ghosh	LIVE	AO	Will it be possible to provide some guidance on A2A in terms of links and online resources and how can I build the knowledge around that myself ?	This program builds on your existing knowledge and focuses on how LLMs enable autonomous systems, where the model acts as the decision and reasoning core of an agent. Also you have completed the Generative AI program - you may feel like the 1st module of the program is an overlap as it covers the similar topics and from module 2 onwards the content is actually focusing on the Agents part.
30	Luis Urso	LIVE	AO	Dr. William, how is the regulatory scenario in Pharma Industry, related to AI products, in special, the challenges on how to escalate to patients ?	There is no single "best" IDE; the right choice depends on the task, scale, and engineering rigor you need. VS Code is strongest for production-oriented workflows: structured projects, debugging, testing, type checking, linting, Git integration, and reproducible environments via venv or conda. Jupyter is best for exploratory analysis, rapid prototyping, and explanation with inline results, but you must manage hidden state, modularize code into packages, and add tests to avoid notebook brittleness. Colab is essentially hosted Jupyter with free or paid GPUs, convenient for quick experiments and sharing, but weaker for environment control, long-term reproducibility, and data governance. For agentic AI systems, prefer VS Code for the core codebase, configuration, evaluation harnesses, and observability, and use notebooks only for exploration and reporting. Next, study environment management, experiment tracking, and eval pipelines, for example conda or uv, MFlow or Weights and Biases, and unit and integration testing.
31	Scott Hardesty		PO	will you be providing this slide deck?	Yes, we can point you to solid starting resources on agent-to-agent (A2A) patterns, including multi-agent orchestration, message protocols, and evaluation. To build knowledge independently, implement a small two-agent system with explicit schemas and a shared state store, add tracing and replay, then stress-test it with adversarial and long-horizon scenarios to observe failure propagation. Next, explore topics like consensus and critique loops, tool arbitration, and communication constraints that reduce cascading errors. Check this: https://github.com/a2aproject/A2A
32	Anonymous Attendee		PO	Could you please provide the calendar for the course, like end date, project deadlines, etc?	In pharma, regulation is risk-based: the higher the impact on patient safety, the stronger the requirements for validation, audit trails, change control, and documented human oversight. For any patient-facing escalation, you typically need clear thresholds for handoff to qualified clinicians, strict limits on autonomous actions, and end-to-end traceability for what the system recommended and why.
33	Anonymous Attendee		AO	how did you make sure that we will use the best & most advanced tools for our projects and which are they?	We will be uploading the slide deck on Olympus after the orientation session.
34	James Schleicher		PO	Can we get the slides presented?	For the schedule of all Mentored Learning Sessions, assessment deadlines, and holiday breaks, please refer to the Learning Schedule available under the Program Overview course. You can access it by following this path: Courses -> Program Overview -> Learning Schedule.

35	Dino Vitale	LIVE	AO	will we cover hot topic of agent evaluation and guardrails (so many new vulnerabilities introduced)	We will touch both evaluation and guardrails in the course.
36	Martin Mugumya	LIVE	AO	Could Agentic AI systems potentially replace Customer relationship management systems?	It may not fully replace human support, but it can reduce reliance on humans for routine requests. Critical, ambiguous, or high-impact cases should still be escalated to human agents for final resolution.
37	Steven Malatji		PO	when did this program start, is 2026 the first group or it started in 2025	The program's first iteration was launched in September 2025.
38	David E solorio	LIVE	AO	With all of the agent platforms coming out, how soon will this tech move from expert built to consumer built? Or will this always require a developer or engineer to implement, or is there a mix?	Several companies are already building custom agents tailored to specific industries, and some are demonstrating strong performance within those narrowly defined domains.
39	Meghan Osburn		PO	Question about the privacy policy: Can you clarify the policy about social login data reuse? This reads broader than just authentication permissions, including the right to store, use for future job applications, and explicitly ongoing reuse. I haven't added my LinkedIn to my profile yet because of this	I would like to assure you that your data and information are completely secure with us. Our LMS, Olympus, is an in-house proprietary system, and any information we request is solely to understand your experience so we can align you with the most suitable mentor and group. Linking your LinkedIn profile is entirely optional. The purpose of linking it is to make it easier for you to attach or share your e-portfolio on LinkedIn. If you prefer not to link your profile, you may simply choose the skip option.
40	Patel Shreyans		AO	Are you going to cover agents can be built in compliance with regulatory framework e.g. in health care industry needs of FDA or ICH GCP?	Yes, we will demonstrate how to design agents with compliance constraints in mind, using guardrails, policies, auditing, and traceability patterns. We will not provide formal FDA or ICH GCP certification guidance, but we will map technical controls to typical regulated-workflow requirements.
41	Ben Miller		PO	Are there going to be many sessions on Sunday? I am less available on Sunday mornings.	Your program manager will help you in aligning a group as per your preferred day.
42	Anand R		AO	What about the cost of developing these models and is it sustainable as a business, see a lot of articles around the AI bubble?	Costs can be managed with a disciplined approach: narrow scope, strong evaluation, caching and retrieval, smaller models when sufficient, and continuous monitoring of spend per task, and we will touch these levers in the course.
43	Robert Peterson		AO	Curious on the staff's perspective on ChatGPT vs. GitHub + Copilot vs. other for Python code acceleration and debugging. GL mentions we can use the platform of our choice (Claude in my case), any upside or downside to either (other than GitHub + Copilot being native to VS Code)?	Copilot is best for fast in-IDE edits and autocomplete in VS Code, while ChatGPT or Claude are better for deeper debugging, refactoring plans, and explaining failures across modules. The main downside of chat tools is weaker project context and more manual workflow, so many teams use Copilot for tight coding loops and a chat model for design review and troubleshooting.
44	Peter Weisz		AO	how did you select the best and most advanced tools for us completing our projects and which are they?	We selected tools based on proven industry impact and engineering maturity, emphasizing reproducibility, integration depth, evaluation support, and governance readiness. The stack we highlight includes RAG components, LangChain and LangGraph, DSPy, SmolAgents, DeepEval, Guardrails, MCP, and RL-oriented methods where appropriate.
45	Salman Sajid	LIVE	AO	How do you see automation and agentic AI together	Automation executes predefined rules, while agentic AI adds the ability to choose and sequence actions under uncertainty by interpreting context with LLMs. The key difference is controlled autonomy: an agent can plan, call tools, and adapt when inputs are messy or the path is not fully specified. In practice, you get the best results by keeping deterministic automation as the backbone and placing agents at decision points, exception handling, and interface layers to unstructured data. This reduces brittleness but introduces new needs for constraints, logging, and evaluation. Next, study orchestration graphs, tool contracts, and human approval gates for high-impact actions.
46	Soumyadip Mukhopadhyay		AO	governance around agentic AI and GenAI solutions and approvals required before deployment often delays the projects.can we get some guidance around effective and efficient governance practices that ensures relative quicker production deployment of these solutions?	Yes, we will discuss governance patterns that reduce deployment friction without weakening controls, by making requirements explicit and testable early. This includes risk-tiering use cases, least-privilege tool access, and clear human-approval gates for high-impact actions.
47	Abhishek Treshan		AO	I work for an top ai financial institution, how much of, the actual of this course relates to using Microsoft ai aka coplit	This course focuses on how modern LLMs enable agents that can plan and act with controlled autonomy. While you will not rebuild Copilot itself, you will understand the core mechanisms behind systems like Copilot and many other AI agents, including tool use, orchestration, and evaluation.
48	Salih Abuelyaman		PO	Do we know what time the mentorship sessions will be? Can we give preference on time?	The time for your mentored learning sessions will be shared by Wed next week. You can connect with your program manager for sharing a preference.
49	Sandip Arvind Chaudhari	LIVE	AO	Can you give any example of use case that failed in agentic AI ? why ? May be explain what should be correct approach ?	A common failure is a customer support "auto-resolution" agent that was allowed to issue refunds or close tickets; it hallucinated policy details and took irreversible actions, creating financial and trust damage. The correct approach is least-privilege tool access, human approval for high-impact steps, and an evaluation suite that stress-tests policy adherence and edge cases before rollout.
50	Avinash Kumar		AO	Is the testing framework around agents built covered in any fashion? Like any traditional engineering lifecycle, how to ensure outcomes are validated to build confidence and trust in solution.	We also cover this aspect by demonstrating how to evaluate agents using a library called DeepEval.
51	Peter Weisz	LIVE	AO	will we be able to produce agents that are functional for our own scope of work or are they pre-defined by the program?	The program will give you the understanding needed to build your own agents tailored to your specific use cases.
52	Pratik Shanu	LIVE	AO	Human do make mistakes but businesses then know who is accountable. Who do you make accountable when an agent makes a mistake? And how to make stakeholders comfortable with that especially in a highly regulated finance domain.	Accountability never sits with the AI, it remains with the business owner and designated human approvers, just like any automated system in finance. Agents should recommend actions, while risky decisions pass through human approval gates with clear audit trails. This preserves traceability, compliance, and stakeholder confidence without relying on AI autonomy.
53	Banketeshvar Narayan	LIVE	AO	Agentic AI is already moving toward maturity. What do you think will be the next wave in AI over the next 1-2 years.	Over the next 1-2 years, the next wave will be multimodal, autonomous systems that can reason, plan, and act across text, vision, audio, and tools moving from task execution to end-to-end workflow ownership.
54	Martin Mugumya		AO	Will the course include building LLMs from scratch? Or will we have to rely on already existing models/APIs?	No, we do not develop LLMs from scratch. Instead, we access pre-trained models through OpenAI's APIs.
55	Dino Vitale		AO	Will we use more LangGraph vs other frameworks as seems like Langchain/Langgraph more used in enterprises?	Here in the Program we cover - SmolAgents, Langgraph/Langchain for building AgenticAI systems, also we cover the overview on CrewAI & Autogen
56	Deepak Kannappan		AO	Do you see any other languages supporting packages similar to python for training and building agents?	Many languages can be used to build agentic systems, but Python is the most common for rapid prototyping and course demonstrations because of its ecosystem of LLM, RAG, and evaluation libraries.
57	Sumit Sharma		PO	Is there any external resources that I can refer during the or before the course	You can study worked examples using SmolAgents, LangChain, and LangGraph, along with a focused set of prompt engineering patterns. https://huggingface.co/docs/smolagents/en/index https://docs.langchain.com/ https://platform.openai.com/docs/guides/prompt-engineering
58	Pranay Lingareddy		AO	What are some helpful resources to keep up with the latest trends in AI.	Track current research through venues like IEEE journals and arXiv, and follow practical implementations and postmortems shared by practitioners on forums such as X and Reddit.
59	Ananda Reddy		AO	Do we touch upon AI governance in this course ?	Yes. There is a dedicated module on governance, guardrails, ethics, and compliance. This includes how to control, monitor, and responsibly deploy agentic AI systems.
60	Ananda Reddy		PO	Can you send me the course brochure for this course ?	Your program manager will share the brochure with you. You can find the details regarding Program Structure in the Program Overview Course> Learner Handbook.
61	Sournava Bhattacharjee		AO	Will this course help us build AI agents to automate regular work in IT operations area?	We will address parts of this domain through our use cases, and the concepts you learn should help you scope, design, and build the specific use case you have in mind.
62	Ankur Vohra		AO	Is it possible to cover a financial services use case project? Like trading risk analysis or portfolio manager agent ?	We will cover some finance-related agent use cases during the course.
63	Juan J. Cerrolaza		AO	It seems that one of the main challenges for the success of many projects and initiatives on Agentic is the problem selecting proper use cases. Do you have any kind of check-list or criteria to identify a good project candidate? What does defines an ideal problem to be solved with agents ai?	A good agent candidate has clear business value, a well-defined success metric, and a workflow that mixes unstructured inputs with structured actions across existing tools. It should be feasible to constrain with policies and permissions, and you must be able to build an evaluation set that represents real cases and failure modes. Avoid problems where errors are catastrophic, requirements are ambiguous, or tool access and data quality are poor. Strong candidates are high-volume, repetitive processes with measurable outcomes, clear escalation paths, and easy human review. Next, formalize this into a small pilot with an ROI model, a gold dataset, and acceptance tests before scaling.
64	Krishna Chaitanya Penumatcha		AO	Is there any course material available on the topic of software architecture? I mean specific to Agentic AI systems	Yes, we will cover architecture patterns specific to agentic AI systems, such as tool and data access layers, orchestration graphs, memory and state management, and evaluation and observability. For deeper study, look at references on LLMops and agent workflow design, plus systems topics like distributed tracing, policy enforcement, and reliability engineering.
65	Vijay Bhimrao Patil	LIVE	AO	Will this course help me to build AI agents for trading who can do the trading automatically without manual intervene ?	Yes, you will learn the core foundations needed to build such systems, so you should be well-prepared to start your own project and iterate with some engineering practice.
66	Raghunandan Ghagarvale		AO	How effective is AI in law domain ? any inherent challenges you see AI not able to solve?	We do not track specific company implementations, but current LLM capabilities can reduce manual effort in legal research, summarizing statutes and prior cases, and drafting and reviewing routine documents under appropriate human oversight.
67	Anonymous Attendee		AO	AI is not new but the growth of AI is exponential now. We always underestimate the exponent - would you agree?	Often yes: when capability scales nonlinearly with compute, data, and integration, our intuitions tend to underestimate the pace, though real-world bottlenecks can still show deployment.
68	Salih Abuelyaman		PO	What's the best way to connect? The WhatsApp link says the number doesn't exist for me.	The support channel details are available under Connect with Us section of Program Overview. You can email us at support.agenticai-jhu@mygreatlearning.com or raise a support request. We will also have your program manager connect with you.
69	Satya Nand Jha		AO	my Firm uses CoPilot as approved platform. Whatever we learn here, will we be able to do those agent creation on Copilot platform?	Copilot is built on large-scale foundation models trained on massive datasets and significant compute, so you will not recreate something at that level within this course. What you will gain is a clear understanding of how such systems work and how to build smaller, well-scoped agents and workflows using similar underlying principles and tooling.

71	Madysen Hamner		AO	what are some job opportunities that would be great to work in with Agentic AI experience	Strong opportunities include AI agent engineer, LLM application engineer, AI systems or platform engineer, and ML engineer focused on RAG, tool integration, evaluation, and observability. You can also target domain roles like automation architect, AI product engineer, or security and governance for AI, where reliability and controls matter.
72	Eszter Csernai		AO	Do you see a danger that delegating more and more skills and even critical thinking to AI agents erodes these skills in humans?	If you treat them as the sole source of truth and stop applying your own reasoning, then yes, skills can atrophy. Used properly, they function as assistants for repetitive tasks, freeing you to focus on higher-level design, critical review, and creative problem solving.
73	Ananda Reddy		AO	If use cursor for my coding assistance - is my code is uploaded into llm cloud ? in other words is my source code is leaked?	Yes, many such tools send code snippets to a cloud service for analysis and completion. Most offer privacy or enterprise modes that limit retention and training use, but the guarantees depend on your plan and settings. Review Cursor's documentation and policies to verify what is transmitted, stored, and eligible for model training.
74	Bheem Singh		AO	It seems AI will replace >80% of workforce on various level of working roles in almost all industries. How will economy sustain if workforce does not have job?	Forecasts of large-scale displacement are uncertain, but history suggests technology shifts job composition more than it eliminates work entirely. The key issue is transition: roles will change, new tasks will emerge, and reskilling plus mobility will determine how disruptive the shift feels. Economic sustainability depends on productivity gains being translated into wages, new demand, and policy mechanisms that manage inequality and labor market shocks. In other words, the question is less "no jobs" and more "different jobs, unevenly distributed over time." You should also examine complementary factors like task automation versus job automation, and the role of regulation and social safety nets.
75	Kolton Starr		AO	Do you think Agentic AI will hurt low-code / no-code platforms like Salesforce or ServiceNow? Since custom apps could be developed so much faster with Agentic AI?	It is already happening: AI-native IDEs such as Cursor are widely used, and many teams now rely on them to generate and refactor substantial portions of their code under standard review and testing practices.
76	James Schleicher		AO	Will this course cover training a deep learning model (like an LLM) or can you point us to a paper or book on this?	No, this program does not specifically cover how to train deep learning models or LLMs. Instead, it focuses on using LLMs to design and build agentic systems for real-world use cases.
77	Omotolani Lewis		AO	Are the example use cases of agentic AI in the field of healthcare / medicine . What are the best applications in these fields and what are the main challenges?	There are many promising applications, but full adoption in clinical settings typically requires strong human-in-the-loop oversight, since agents should not take autonomous actions that could directly affect patient health. In the nearer term, agentic AI is more likely to see broader use in medical research, literature synthesis, trial operations support, and discovery workflows, where outputs can be reviewed and validated.
78	Shekhar Avasarala		AO	Dr. William, regarding failures: As we move from single agents to 'swarms' or multi-agent systems, have you seen cases of 'cascading failures'—where one agent's small hallucination gets amplified by the next agent in the chain? How do we debug that effectively?	Yes, cascading failures can occur in multi-agent or swarm setups when one agent's mistake propagates through downstream agents. There are established multi-agent design patterns to synchronize and constrain agent interactions, and we will touch on these in the course, along with practical debugging approaches like tracing, replay, and stepwise validation.
79	Ardendu Mishra		PO	Will GL team create a whatsapp collaboration group for collaboration between members through out this program ? previous learning sessions we had similar groups created by GL team, and was very helpful.	We will be creating whatsapp groups before your first mentored learning session next week. We will be sharing the group joining link via email.
80	Anonymous Attendee		AO	One of the critical problem enterprises are trying to solve is around interfaces between Agentic AI solution and various enterprise systems like ERP, CRM etc. Any design patterns and approach going to be covered?	Yes, we will cover key integration patterns for connecting agents to enterprise systems like ERP and CRM, including standardized tool interfaces and access control. We will also touch on approaches like Model Context Protocol (MCP) to connect legacy data and applications to LLM-based agents.
81	Deepak Kannappan		AO	How hallucinations can be avoided for the agents being built?	There are established prompt engineering principles that help make model behavior more reliable and repeatable. We will cover these techniques as we progress through the course.
82	Thomas R Ladue		AO	Is there a parallel relationship between AI and Moores Law...is AI going to accelerate Moores Law in the grand scheme of IT?	There is a partial parallel: AI progress is not governed by Moore's Law, but it can indirectly accelerate compute scaling by helping chip teams optimize architecture, layout, verification, and manufacturing yield, which improves performance per watt and time to design. That said, this is more "AI-assisted scaling" than a restart of classic transistor-density doubling. Long-term limits still matter, especially the memory wall and data-movement costs, along with power and thermal constraints. So AI contributes to faster iteration and better designs, but it does not remove the fundamental bottlenecks that constrain Moore-style scaling. Next, look into compute trends like Dennard scaling end, HBM and memory hierarchy, and hardware/software co-design.
83	Julie Regairaz		AO	Would it be possible to address opportunities of Agentic AI for emerging countries' (mainly on the African continent) challenges in order to reduce the digital divide?	The use cases we demonstrate are intended to illustrate what is feasible with current agentic AI methods. If you have an idea you want to brainstorm, you can use the course community and support channels to get feedback and refine it.
84	Keld Hansen		AO	If a large company decide that all support is based on AI, how will a user then be responsible for an error if that is the only choice to get support and help?	If AI-only support is the sole channel, accountability must remain with the organization through clear audit trails, human escalation paths, and defined liability and remediation processes, not with the end user. We will touch Human-in-the-loop regarding this.
85	Babu Shobin Jose		AO	What are the biggest areas where Agentic AI can make an impact in year 2026?	In 2026, the biggest impact will be in high-volume knowledge work where agents can reliably execute end-to-end workflows: customer support resolution, IT and security ops, finance and procurement, and software delivery with strong testing and reviews. The limiting factors will be integration depth, evaluation and monitoring, and governance, so the largest gains will come where tools, data, and policies are already well-structured.
86	Sahitya Gollapudi		AO	What resources do we have access to? For the Projects, do we get access to use any LLMs as part of the course?	Yes, throughout the program, we will provide access to OpenAI LLMs for your use. The same access will be used across all three projects, and you do not need to purchase any separate LLM subscription. Everything required will be provided by Great Learning.
87	Jayesh Ambadas Sangpal		AO	I see loads of Duplications of work as the Agentic ecosystem kick off. Any thoughts on Reusability frameworks ?	Yes, early duplication is normal, so you need a reuse layer: shared tool APIs (MCP or an internal SDK), common schemas, and a central catalog of approved tools and agent components. Also version prompts, policies, and eval tests so teams can reuse proven pieces instead of rebuilding them.
88	Ashwin Rajput		AO	There is talk about emerging topics like Physical AI which will get AI in to physical world, any thoughts on that ?	We do not explicitly cover Physical AI in this program. However, we do cover interaction and Embodied AI. Embodied AI focuses on integrating artificial intelligence with physical systems - such as robots or vehicles - enabling them to perceive, learn, and act autonomously in real-world environments.
89	Ardendu Mishra		AO	is it mandatory to use python for this program or other languages like java can be used for the handson ?	Yes, our current projects as well as all hands-on sessions are designed using Python.
90	Meghan Osborn		PO	Is it technically possible to share transcripts? I am happy to explain the request offline	We do not have the transcripts available for the recorded lectures or live sessions on your end. However, you may reach out to your Program Manager, who can assist you with transcripts for each Mentored Learning Session. Regarding the recorded lectures and content, we have in-built LLM model, Glaiide aligned with each lecture. You can use Glaiide to ask questions or get a summary of any Video at your convenience.
91	Kolton Starr		AO	Do AI Agents make have a place in the context of very deterministic workflows?	Yes, they do. With the capability of autonomy added due to LLMs, it can understand things in that workflow that was not possible earlier.
93	Leonardo Netto		AO	From a technical efficiency perspective (ignoring privacy/regulation), how do you decide when to create a new AI agent versus extending an existing one? In a large corporation, is a single general-purpose agent for an entire corporation realistic, or is it more effective to scope agents per department, function, or role? What factors usually define that boundary?	Create a new agent when the task has distinct objectives, tools, data, and failure modes that would make an existing agent's policy brittle, hard to evaluate, or expensive to run: extend an agent when you can add capabilities without degrading reliability, latency, or test coverage. A single corporate "general agent" is rarely realistic beyond a thin routing layer, because different domains need different permissions, ontologies, and guardrails. In practice you use a hub and spoke pattern: a router plus multiple specialized agents scoped by workflow. Boundaries are usually defined by tool access and least privilege, data ownership, domain terminology, required accuracy and auditability, and clear evaluation suites per role. Next, explore agent routing, capability registries, and multi-tenant policy enforcement.
94	Ana Saxena		PO	just as a follow on, we can't see the content aearly and get ahead on self reading. I understand the projects and quizzes are weekly but would be good to give the opportunity to read ahead.	As this is a structure learning journey, the course content is available to everyone on Olympus as per the schedule.
95	Lucia Izquierdo Rajadel		PO	Hi, in addition to the mentor live sessions where you can raise questions for the content delivered during the week, how do you manage any outstanding question you may have for topics covered in previous weeks?	In addition to the mentor-led live sessions, you can raise questions related to topics from previous weeks by submitting a Support Request. You can do this by clicking on the question mark (?) icon at the top-right corner of the learning platform. Our academic team review these requests and respond accordingly. You may also revisit session recordings and bring up specific doubts during upcoming live sessions for further clarification. Thank You !
98	Santosh Kumar		PO	Will all the mentored learning sessions for an individual will be at the same time for entire program	You will be assigned to one group for your mentored learning sessions.
99	Miguel Tablado Leon		PO	The first mentoring session is tomorrow, why it says the weekend of 17th?	The Week 1 mandatory Mentored Learning Session is scheduled for the weekend of 17th-18th January. The session on 11th January was a Week 0: Pre-work session, which was optional for you.
100	Anonymous Attendee		AO	When to use google colab vs VS Code?	You can use either option for the projects as well as the MLS sessions. Each option has its own unique features, along with pros and cons. If you prefer working on your local system, you can install VS Code and Python and use them throughout the program. Alternatively, if you want to avoid the hassle of installing software and continue directly on a cloud platform, you can use Google Colab. With Google Colab, you don't need to manage Python libraries individually and can execute the notebooks directly in the cloud.
101	Michael Shea		PO	will the weekend live sessions be at the same time/day?	The weekend live Mentored Learning Session schedule will remain the same throughout the program, unless there is a change in a specific week due to emergency at mentor's end, or overlap with the monthly masterclasses.
103	Narciso Tavora		PO	do we have another chance if we fail de deadline assessment?	You will get one attempt for each graded quiz. In case of any technical issues, please contact your Program Manager for an additional attempt.

104	Zvonko Pino		PO	If my time zone changes during the course, will everything adjust automatically?	If your time zone changes during the course, please update it under Settings. This will ensure that the schedule on Olympus and all session timings adjust automatically to reflect your current time zone, so you always see the correct session times, deadlines, and notifications.
105	Babu Shobin Jose		PO	Do we have option to retake graded tests?	You will get one attempt for each graded quiz. In case of any technical issues, please contact your Program Manager for an additional attempt.
106	Shariq Ahmed		PO	Is the first graded quiz due on January 19th at 3am? Could you please specify the time zone? Is it Eastern Time Zone?	The first graded quiz is due on January 19th at 1:25 PM EST / 11:25 PM IST.
107	Meghan Osborn		PO	Thanks for this overview! FYI - I flagged a concern in the Q&A earlier about logging in with social credentials and the privacy policy terms	I would like to assure you that your data and information are completely secure with us. Our LMS, Olympus, is an in-house proprietary system, and any information we request is solely to understand your experience so we can align you with the most suitable mentor and group. Linking your LinkedIn profile is entirely optional. The purpose of linking it is to make it easier for you to attach or share your e-portfolio on LinkedIn. If you prefer not to link your profile, you may simply choose the skip option.
108	Lucia Izquierdo Rajadel		PO	Hi, what are the requirements (% of attendance, material delivered, etc) to obtain the certificate at the end of the program ?	To be eligible for the certificate please ensure that you score minimum 60% in each module
109	Ana Saxena		PO	can we get ahead, or would the sessions only be unlocked week by week?	The program will follow a weekly learning structure, so every week you will cover 1 topic. The content is also shared weekly.